

Multi armed Bandits

Chuck Chan

April 3, 2021

Problem

- Conditions
 - Fixed set of limited resources and time
 - Resources must be allocated between competing choices to maximize expected gains
 - Each choice has properties that are partially known at the time of allocation and becomes better understood as time passes or by allocating resources
 - Assumes the distribution for rewards from choices are fixed but unknown
- Classic Example
 - Gambler deciding which slot machine to play and how many times to play each, which order to play them, and if they should continue or switch machines.

Explore & Exploit tradeoff

- A tradeoff is created between time spent on Exploring and Exploiting the choices
 - Exploration is the time spent gathering information on the possible choices.
 - Exploitation is spending the time on the choice that is expected to maximize the gains.
- Regret (ρ) is the difference between gains from the strategy taken and maximum gains given complete information. It is used to evaluate how well a strategy performs.
- Some algorithms can be stuck on suboptimal solutions like local means

Multi Armed bandit Applications

- Clinical Trials
 - Adaptive allocation strategy to improve data collection by allocate more samples for exploring promising treatments
- Portfolio management
 - Making portfolio choices by exploiting correlations among arms to balance risk and amount of passive and active investments
- Dynamic pricing
 - Real-time pricing for product balancing demand with maximizing profits
- Recommender Systems
 - Recommend relevant items to users based on contextual information
- Social Media Marketing
 - Maximize user awareness of a product through social networks by selecting seed users for exposure
- Information Retrieval
 - Determines selection process for information returned by search
- Anomaly detection
 - Identify deviating nodes that are used for predicting fraudulent transactions that are given to a human investigator

Solutions

- ϵ Greedy
 - Assigns probability to balance exploration / exploitation
- Softmax / Boltzmann Exploration
 - Algorithm converges on an efficient arm
- Pursuit Algorithm
 - Modifies probability of choosing each arm based on rewards
- Reinforcement Comparison
 - Calculates a preference for each arm and compares an mean rewards
- Upper Confidence bound family of algorithms
 - Construct and updates a mean and confidence interval for the reward of each arm
- Gittens Index
 - Maintains a value function for each arm
- Thompson Sampling
 - Maintains a mean and standard deviation for each arm and selects based on prior probabilities

ϵ Greedy Strategies

- Strategy to balance explore and exploit options
 - ϵ is the probability of choosing explore and $1 - \epsilon$ is exploit
 - Each round an arm is selected based on a probability is generated randomly of explore and exploit
 - During exploration each non-best arms are all considered equivalently
- Represented as:

$$p_i(t+1) = \begin{cases} 1 - \epsilon + \epsilon/k \\ \epsilon/k \end{cases}$$

Probability of choosing arm i at time $t+1$

Probability the best arm gets played
Corresponding to $i = \operatorname{argmax}_{j=1,\dots,K} \widehat{\mu}_j(t)$

Probability a random arm getting selected

k is the number of different arms

Softmax / Boltzmann Exploration

- Select arms based on rewards from a probability from the Boltzmann distribution
- Arms with larger empirical means have a higher chance of being picked
- Hyperparameter τ controls the randomness of the choice
 - $\tau = 0$ means the algorithm becomes greedy picking the largest empirical mean
 - When τ is large, arms are selected uniformly at random
- Represented as:

$$p_i(t+1) = \frac{e^{\widehat{\mu}_i(t)/\tau}}{\sum_{j=1}^k e^{\widehat{\mu}_j(t)/\tau}}, i = 1, \dots, n$$

Diagram annotations:

- A green arrow points from the text "average reward for arm i" to the term $\widehat{\mu}_i(t)$ in the numerator.
- An orange arrow points from the text "average reward for all arms" to the term $\widehat{\mu}_j(t)$ in the denominator.

Probability of choosing arm i at time t+1

Pursuit Algorithm

- Starts with uniform probabilities assigned to each arm $p_i(0) = 1/k$ and a learning rate $\beta \in (0,1)$
- Probabilities are updated after each turn t
- Represented as:

$$p_i(t+1) = \begin{cases} p_i(t) + \beta(1 - p_i(t)) \\ p_i(t) + \beta(0 - p_i(t)) \end{cases}$$

Update if $i = \operatorname{argmax}_j \widehat{\mu}_j(t)$

Update of arm i
otherwise

Learning rate between (0,1)

Reinforcement Comparison

- Maintain an distribution over all actions that is updated
- Maintain an average expected reward $\bar{r}(t)$
- Probability of selecting an arm is computed by comparing empirical mean with the average expected rewards
 - Probability is increased if it is above the average expected rewards
 - Decreased if equal or lower
- Probability of selecting an arm i :

$$p_i(t) = \frac{e^{\pi_i(t)}}{\sum_{j=1}^k e^{\pi_j(t)}}$$

Preference for each arm i

average preference for arm j

- Preference update if arm $j(t)$ is played and reward $r(t)$ is received:

$$\pi_{j(t)}(t+1) = \pi_{j(t)}(t) + \beta(r(t) - \bar{r}(t))$$

Average expected rewards

- Mean rewards update at every turn:

$$\bar{r}(t+1) = (1 - \alpha)\bar{r}(t) + \alpha r(t)$$

Learning rates between 0,1

Reward at time t

Upper Confidence Bounds algorithm family

- construct a confidence interval of what each arm's true performance might be
 - Factors the uncertainty caused by variance using empirical means
 - Limited series of pulls tracking times each arm is played
- Algorithm assumes each arm will perform as well as its upper confidence bound selecting the arm with the highest one
- Algorithm variants
 - UCB-1
 - Tuned UCB-1 (includes variance)
 - Bayes UCB (includes prior distribution information)

Upper Confidence Bounds 1

- maintains the number of times each arm is played and empirical means
- Use Hoeffding's inequality to assign upper bound to an arm's mean reward
- Initially each arm is played once
- Followed by a greedy algorithm:

$$j(t) = \arg \max_{i=1, \dots, k} \left(\hat{\mu}_i + \sqrt{\frac{2 \ln t}{n_i}} \right)$$

confidence interval for the average reward

Number of times arm i is played

average reward for arm i

Upper Confidence Bounds 1

- Smaller n_i the larger the Confidence Intervals and less accurate mean rewards, Larger results in smaller confidence and more accurate mean rewards
- As t increases the algorithm converges to the optimal action
- Regret calculation

$$8 \sum_{i: \mu_i < \mu^*} \frac{\ln(t)}{\mu^* - \mu_i} + \left(1 + \frac{\pi^2}{3}\right) \sum_{i=1}^k \mu^* - \mu_i$$

- Regret is $O(\log(n))$ bounded growth, optimizing regret to a multiplicative constant

Tuned UCB1 (Audibert, et al 2009)

- Performs better than UCB
- Includes variance of each arm
- Equation to pick arms at turn t:

$$j(t) = \arg \max_{i=1, \dots, k} \left(\hat{\mu}_i + \sqrt{\frac{\ln t}{n_i} \min \left(\frac{1}{4}, \hat{\sigma}_i(t) + \frac{2 \ln t}{n_i} \right)} \right)$$

Highest variance for
reward distribution of
Bernoulli random variable
($p=0.5$)²

Variance calculated by
sum of squares of
reward

average reward for arm i

Number of times arm i is played

Bayesian UCB

- Uses Priors about distributions in determining confidence intervals
- Assume rewards normally distributed
- Algorithm:

The diagram shows the formula for selecting the best arm $j(t)$ at time t using the Bayesian UCB algorithm. The formula is
$$j(t) = \arg \max_{i=1, \dots, k} \left(\hat{\mu}_i + \frac{c \sigma_i(t)}{\sqrt{n_i}} \right)$$
 where $\hat{\mu}_i$ is the average reward for arm i , c is a hyperparameter, $\sigma_i(t)$ is the standard deviation of rewards of arm i , and n_i is the number of times arm i is played. The annotations are as follows:

- Hyperparameter determining the size of Confidence interval** (purple text, arrow pointing to c)
- Standard deviation of rewards of arm i** (red text, arrow pointing to $\sigma_i(t)$)
- average reward for arm i** (green text, arrow pointing to $\hat{\mu}_i$)
- Number of times arm i is played** (blue text, arrow pointing to n_i)

Thompson Sampling

- also known as posterior sampling & probability matching
- Aims to quickly identify arm with highest mean, by starting with a posterior belief for the mean and adjusting the mean for the arm played
- Algorithm converges to the optimal solution.
- Chooses arm with posterior probability of being the best arm
 - Not choosing the one most likely to be effective
 - Gives benefit of doubt to less explored choices
- Accounts for uncertainty about means, higher uncertainty ensured a less explored arm is pulled
- The variance in the posterior enables exploration

Thompson sampling for Bernoulli rewards

- Sampling distribution with Bernoulli rewards
- Start with Beta priors using $Beta(1,1)$ for each arm
- $Beta(\alpha, \beta)$ priors are updated at round t where:
 - Successes updates $Beta(\alpha+1, \beta)$ if 1 is observed
 - Failures updates $Beta(\alpha, \beta+1)$ if 0 is observed
- Arm i is selected that maximizes rewards calculated by $\hat{\theta}_k = \frac{\alpha_k}{\alpha_k + \beta_k}$
- Beta posteriors are updated.

Thompson Sampling with Gaussian Priors

- Rewards for arm i is $N(\mu_i, 1)$
- Starting priors set $N(0,1)$
- Sample $\tilde{\mu}_{i,t}$ from posterior mean for $N\left(\hat{\mu}_{i,t}, \frac{1}{n_{i,t}+1}\right)$ arm i
- Arm i is selected that maximizes $\tilde{\mu}_{i,t}$
- Update the empirical mean $\hat{\mu}_{i,t}$ for i_t

Empirical mean



Number of observations



Gittins index

- Measures reward that can be achieved in a stochastic process using dynamic allocation indexes
- Used for Bayesian MAB problems where
 - Each arm has a reward distribution that has an known prior
 - The arms are independent
 - The objective is to maximize the expected discount reward
- Prior probabilities are updated each time a decision is made with posterior probability
- An single arm is represented as a Markov chain and reward function
- Policy is a function that determines the decision of which arm to chose next.

Gittens Index Calculation and Policy

- A Gittens Index (GI) is calculated for each arm and the reward is allocated as stochastic process of success and failure with an unknown probability of success
- A policy is implemented and depends on the relative accuracy of the GI limited by:
 - decisions involving arms with similar GI values being rare
 - Cost of a suboptimal action is small when the GI is close to the optimal arm, giving a boundary for the loss rewards of the policy in terms of GI where the GI accuracy limits the losses and rewards.
- Each time an arm is selected the GI for that arm is updated.
- Some policies may reduce / retire arms

Calibration method for GI calculation

- Value function for single arm(risk arm):

$$V(\Sigma, n, \gamma, \lambda) = \max \left\{ \underbrace{\frac{\Sigma}{n} - \lambda + \gamma E_Y [V(\Sigma + Y, n + 1, \gamma, \lambda)]}_{\text{Value of the risk arm}}; \underbrace{0}_{\text{Value of the safe arm reduced to 0}} \right\}$$

- The fix reward arm λ minimizes the value function to close to zero and is found by recursion of a calibration algorithm given some initial bounds

Variables

λ an arm of known fixed reward

γ discount factor

Y random variable of the risk arm state with state (Σ, n)

Σ Bayes sum of rewards of the risk arm

n Bayesian number of observations of the risk arm

Bernoulli MAB GI calculation

- Reward of one or zero
- Value function for discrete single arm of a disc

$$V(\Sigma, n, \gamma, \lambda) = \max \left\{ \frac{\Sigma}{n} - \lambda + \gamma \left[\frac{\Sigma}{n} V(\Sigma + Y, n + 1, \gamma, \lambda) + \left(1 - \frac{\Sigma}{n} \right) V(\Sigma, n, \gamma, \lambda) \right]; 0 \right\}$$

Variables

λ an arm of known fixed reward

γ discount factor

Y random variable of the risk arm state with state (Σ, n)

Σ Bayes sum of rewards of the risk arm

n Bayesian number of observations of the risk arm

Algorithms for the multi-armed bandit problem (2000)

- Tested multiple MAB solutions against different bandit characteristics
- Algorithms tested
 - ϵ Greedy
 - Boltzmann Exploration
 - Upper Confidence Bounds
 - Reinforcement Comparisons
 - Pursuit Algorithms
- Characteristics for MABs
 - Different distributions
 - Number of arms
 - Reward variances
- Tested algorithms in Medical Clinical trial settings

Algorithms for the multi-armed bandit problem (2000)

- Simple heuristic algorithms like Boltzmann exploration outperformed advanced algorithms
 - Both ϵ Greedy and Boltzmann Exploration had versions that had decreasing hyperparameters ϵ and τ but were found to have no advantage by Vermorel and Mohri (2005)
- Performance of algorithms affected by number of arms and reward variances
- Algorithms vary depending on Bandit instances
 - UCB family is good for bandits with small number of arms and high reward variance, but degrades as number of arms increase

Adversarial Bandits

- Multi-arm Bandit variant where an adversary determines the outcome of based on the arm that is selected
- Goal is to minimize negative outcomes
- Examples
 - iterated prisoner's dilemma
 - Clinical trials
- Solutions include EXP3 algorithm
 - Weights are adjusted for favorable arms

Exponential weight for Exploration & Exploitation (EXP3)

- Weight are adjusted over time to favor good arms and reduced for less promising arms
- Algorithm initialized a vector of weights each at 1
- Setup:

$$p_i(t) = (1 - \gamma) \frac{w_i(t)}{\sum_{j=1}^k w_j(t)} + \frac{\gamma}{k}$$

Diagram annotations:

- Sum weight of all arms (points to $\sum_{j=1}^k w_j(t)$)
- Weight of arm i (points to $w_i(t)$)
- Hyperparameter for exploration of arms uniformly at random where 1= pure exploration (points to γ)
- k is the number of different arms (points to k)

- Draw arms i_t randomly according to p_i, \dots, p_k and observe rewards $x_{i_t} \in [0,1]$
- Define estimated reward $\hat{x}_j(t) = \frac{x_j(t)}{p_j(t)}$ for $j = i_t$, otherwise 0 for all j
- Update weights $w_j(t + 1) = w_j(t)e^{\gamma \hat{x}_j(t)/k}$

Contextual Bandits

- Multi-armed bandit variant where choosing an arm is dependent on a set of contextual information. The reward is observed for the chosen action
- Goal:
 - Balance explore and exploit behavior
 - Use context as effectively as possible, generalize across context
 - Handle selection bias from skewed data of exploring while exploiting
- Characteristics
 - K arms
 - d-dimensional feature vector $x_{i,t}$ for every arm i , at time t
 - Linear parametric model with parameter θ and expected rewards for arm i , at time $t : x_{i,t}\theta$
 - Optimal arm depends on context $x_t^* = \underset{x_{i,t}}{\operatorname{argmax}} x_{i,t}\theta$

Contextual Bandits

- Determine a policy to choose action based on context
 - Policies decide prior to learning
 - Associating context with the best action is treated as a classification problem
- Examples
 - News / Advertising from a website that is served to users with a set of browsing history
 - Medical doctor assigning a treatment with the most favorable outcome based on symptoms
 - Clinical trial with patient data
 - Recommender systems for media with user preferences

Contextual bandit solutions

- Follow the Leader (FTL)
 - Find the empirically best policy and use it to select the best action
- Thompson Sampling
 - Use least square estimator for rewards, the state transition is learned by comparing against a fixed policy.
 - Transition probabilities for each state and action is stored in a Dirichlet posterior
- EXP4
 - Modified EXP3 algorithm with experts advice influencing probability
- LinUCB
 - Modified UCB algorithm that uses a data matrix of features for context, a linear function with hyperparameters, and a vector corresponding to feedback

More Bandit variations

- Lipschitz Bandit, where an algorithm has information on the similarity between arms.
 - Used for dynamic pricing problems
- Bandits with Knapsack, bandit problems with a supply constraints
 - Used in dynamic pricing problems with limited supply
- Combinatorial bandit, where a set of k arms are selected
 - Used for online advertising and social influence maximation.
- Adversarial Bandits with corruption, where the adversary knows the actor's choice and has the ability to corrupt the reward of the selected arm.

The End